

Conclusions

The **linear model** achieved good results and is the most interpretable model.

The **gradient boosted decision tree** model was able to capture some nonlinear relationships between the predictors achieving a mean squared error 76% smaller than the one from the linear model.

The gradient boosting model achieves good prediction power without renouncing to interpretation.

Possible improvements:

- Result interpretation: further analysis + domain knowledge.
- Parameter running: further hyperparameter optimization.
- Modelling: explore other regression techniques
- Dataset: more features

